

Database & Python ETL

Provision one database of your choosing (SQL, NoSQL, Graph). Write a python ETL that ingests the provided data, transforms it in some way, and loads it into the database. This should be reproducible code with documentation. (Terraform / CloudFormation / Ansible, docker-compose etc.).

Introduction:

I have chosen the **Databricks Platform** for building the ETL, as The Databricks Lakehouse Platform combines the best elements of **data lakes** and **data warehouses**. It empowers data scientists, data engineers and data analysts with a simple collaborative environment to run interactive and scheduled data analysis workloads.

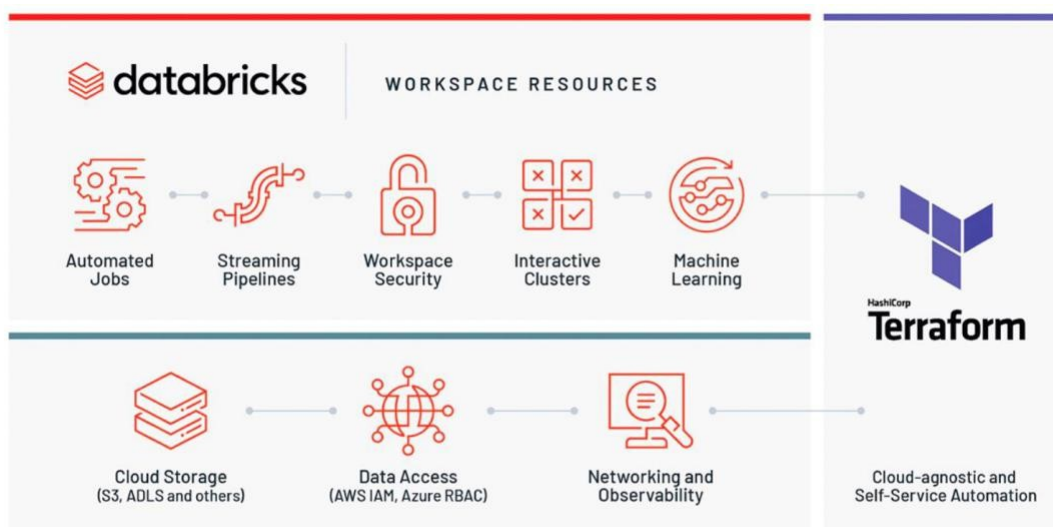
Databricks Terraform provider is to support all Databricks REST APIs, supporting automation of the most complicated aspects of deploying and managing your data platforms. We can deploy and manage clusters and jobs and to configure data access.

Objective:

Goal is to ingest the data, transform it and load it into database tables. Automate complicated aspects of deploying and managing your data platforms using Terraform provider.

Architecture:

Build ETL in databricks notebook, created database in databricks lake house. Used Terraform to automate deployment



Terraform supports management of all Databricks resources and underlying cloud infrastructure

Implementation of the use case

Step 1: Create Azure Databricks Workspace

Step 2: Integrate with GitHub

Step 3: Add repository in Databricks

Step 4: Upload csv file into Databricks Repo

Step 5: Create Python notebook

Step 6: Python code for the below tasks

Σ Data Extraction

- Define the widgets to enter the values of the parameters
- Assign parameter values to variables
- Read the csv file and create pandas dataframe

Σ Data Transformation

- Convert column datatype to int
- Convert pandas dataframe to spark dataframe

Σ Data Loading

- Define source, database name and table name
- Database Creation
- Write/Load the data from dataframe to table

Σ Query the data from the table

Step 7: Create Databricks Job to automate the task

Step 8: Create Job

Step 9: Terraform configuration

Structure of the project

Microsoft Azure | databricks | Search | + P

Data Science & En... ▾

+ New

Workspace

Repos ← Repository

Recents

Data ← Database tables are stored

Compute ← Clusters/ Compute resources are created

Workflows ← Data pipelines/ Jobs are created and managed

Repos

etl-pipeline main ▾

- data-load-transform ▾
- README.md ▾
- SRDataEngineerChallenge_DATA... ▾

Python Notebook to ingest, transform and load into database table

Dataset in csv format

Step 1: Create Azure Databricks Workspace

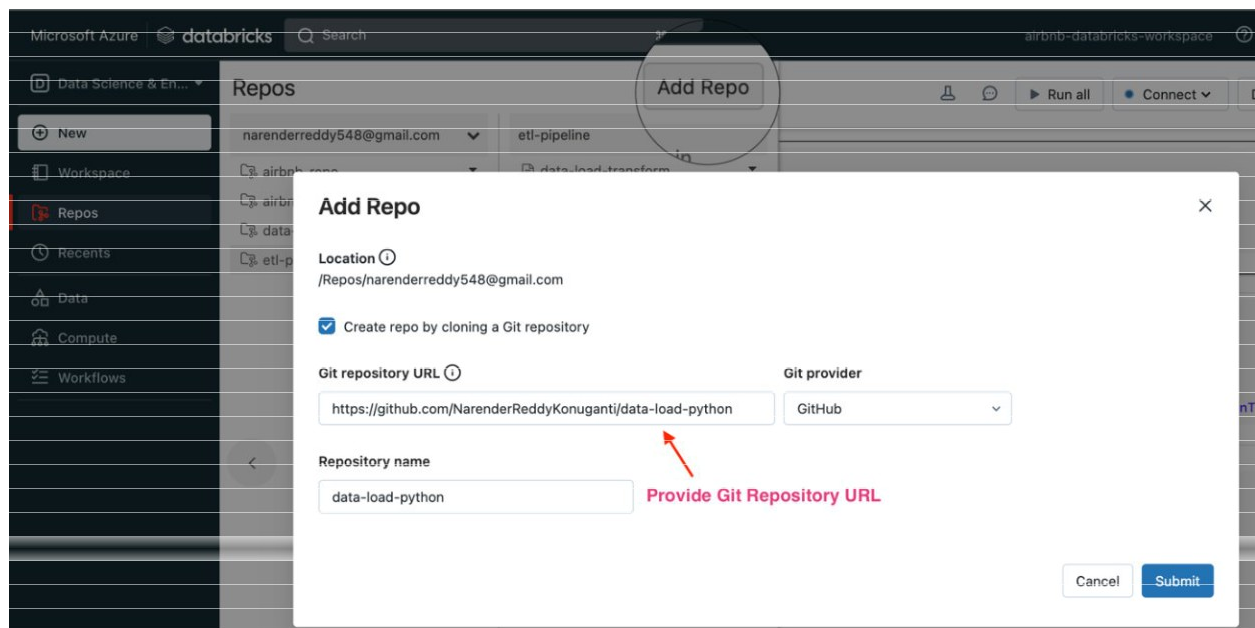
Follow the steps in <https://learn.microsoft.com/en-us/azure/databricks/getting-started/> to create Azure Workspace

Step 2: Integrate with GitHub

Integrate GIT with Databricks using access token, click here for steps [Git-Integration](#)

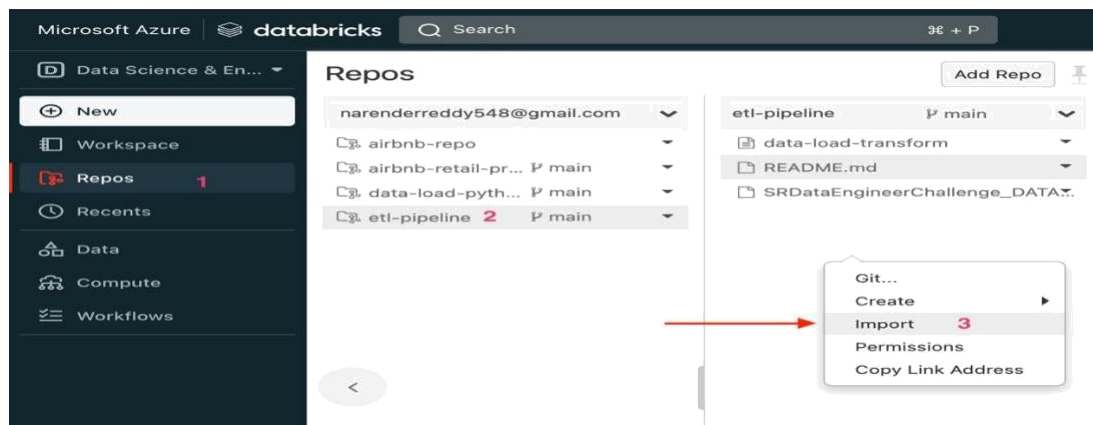
Step 3: Add repository in Databricks

Navigate to workspace ↗ Repos ↗ Add Repo ↗ Provide Git repository URL



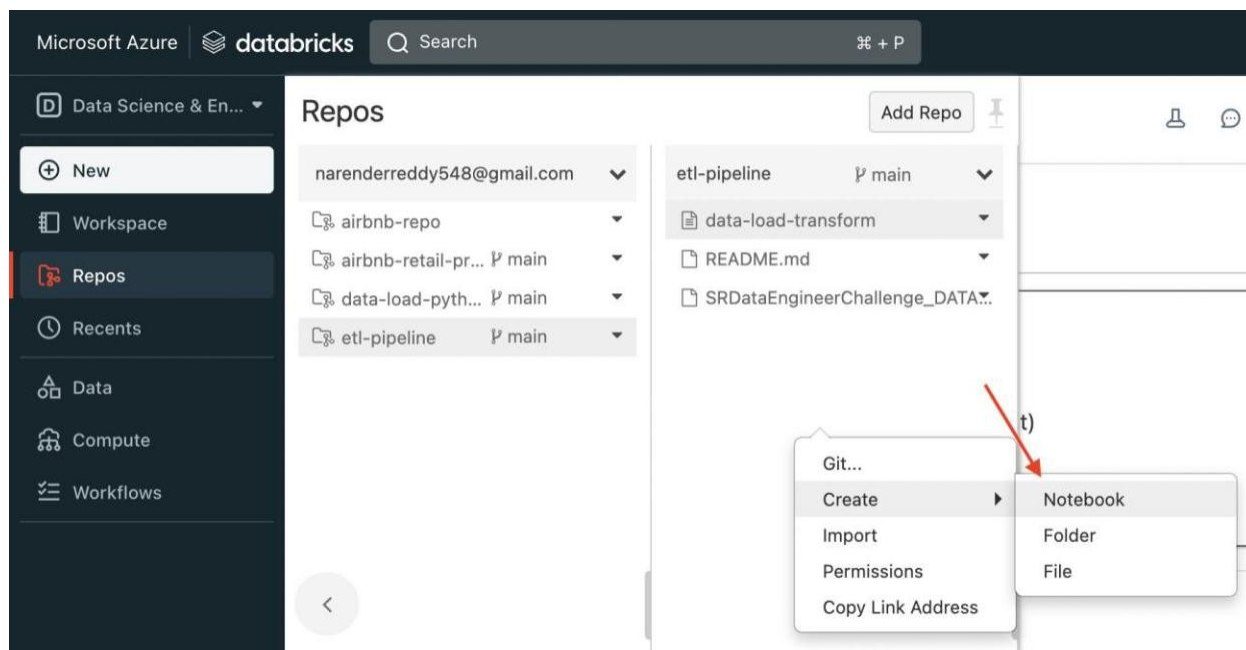
Step 4: Upload csv file into Databricks Repo

Navigate to the etl-pipeline repo and right click ↗ import ↗ upload csv file




Step 5: Create Python notebook

Navigate to the etl-pipeline repo and right click **create notebook**



Create a Cluster:

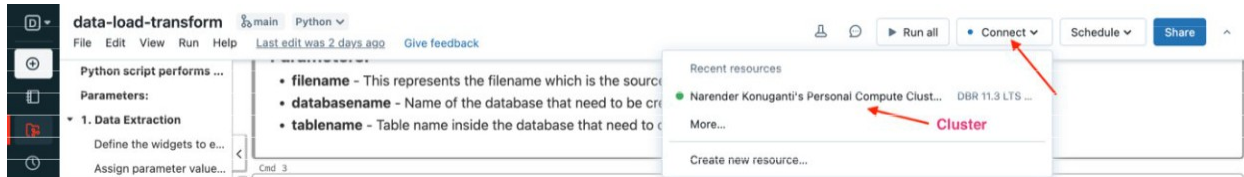
To do exploratory data analysis and data engineering, create a cluster to provide the compute resources needed to execute commands.

- Σ Click  **Compute** in the sidebar.
- Σ On the Compute page, click **Create Cluster**. This opens the New Cluster page.
- Σ Specify a unique name for the cluster, leave the remaining values in their default state, and click **Create Cluster**.

Attach notebook to cluster:

Navigate to the etl-pipeline repo and click on notebook “data-load-transform”

Click on connect and click cluster



Step 6: Python code for the below tasks

Python script performs the below tasks:

- Ingests the data from csv file and creates dataframe
- Transforms the data (Change "id" datatype column to int)
- Creates database in lakehouse
- Creates database tables and loads data

Parameters:

- **filename** - This represents the filename which is the source datasource
- **databasename** - Name of the database that need to be created for storing the data
- **tablename** - Table name inside the database that need to be created for loading the data from dataframe

Data Extraction

- Σ Define the widgets to enter the values of the parameters
 - **Widgets** - Main purpose of widgets is to **parameterize** notebooks. These are used to define the parameters for the notebook.
 - When we **parameterize**, code can be directly deployed in production and run with updated parameter values as per PROD.
 - There are 3 parameters used which is described in above screenshot.
- Σ Assigned parameter values to variables
- Σ Read the csv file and created pandas dataframe

1. Data Extraction

Define the widgets to enter the values of the parameters

```
#dbutils.widgets.remove("TempTableName")
dbutils.widgets.text("filename", "SRDataEngineerChallenge DATASET.csv", "Enter file name in csv format")
dbutils.widgets.text("databasename", "sflatabase", "Enter database name")
dbutils.widgets.text("tablename", "sfltable", "Enter Table name")
```

Assign parameter values to variables

```
import pandas as pd
filename = dbutils.widgets.get("filename")
databasename = dbutils.widgets.get("databasename")
tablename = dbutils.widgets.get("tablename")
```

Read the csv file and create pandas dataframe

```
# Read all as String
hrdata = pd.read_csv(filename, converters={i: str for i in range(100)})
```

Data Transformation

- Σ Converted column datatype to int
 - Transformation - Changed data type of column "id" from **string** to **int**.

- Σ Converted pandas dataframe to spark dataframe

2. Data Transformation

Convert column datatype to int

```
hrdata['id'] = hrdata['id'].str.replace("\$|,", "").astype(int)
```

```
:1: FutureWarning: The default value of regex will change from True to False in a future version.
hrdata['id'] = hrdata['id'].str.replace("\$|,", "").astype(int)
```

Convert pandas dataframe to spark dataframe

```
# From pandas to DataFrame
df_hrdata = sqlContext.createDataFrame(hrdata)
```

Data Loading

- Σ Define source, database name and table name
- Σ Database Creation
- Σ Write/Load the data from dataframe to table

- ❖ In the first step, we have defined the source where the database tables are to be created
- ❖ Specified **database** name and **table name**
- ❖ Created database in databricks lake house
- ❖ Created table inside the database
- ❖ Load the transformed data from dataframe to **database table**

3. Data Loading

Define source, database name and table name

```
username = spark.sql("SELECT regexp_replace(current_user(), '[^a-zA-Z0-9]', '_').first()[0]")
source = f"dbfs:/user/{username}/copy-into-demo"
spark.sql(f"SET c.username='{username}'")
spark.sql(f"SET c.databasename='{databasename}'")
spark.sql(f"SET c.source='{source}'")
```

Out[72]: DataFrame[key: string, value: string]

Database Creation

- Drop database if existing with same name
- Create database
- Define to use the created database

```
spark.sql("DROP DATABASE IF EXISTS ${c.databasename} CASCADE")
spark.sql("CREATE DATABASE ${c.databasename}")
spark.sql("USE ${c.databasename}")

dbutils.fs.rm(source, True)
```

Out[73]: False

Write/Load the data from dataframe to table

- Drop the table if existing with the same name

```
# Write the data to a table.
spark.sql("DROP TABLE IF EXISTS " + tablename)
df_hrdata.write.saveAsTable(tablename)
```

Query the data from the table

Below is result of the data retrieved from table.

Query the data from the table

```
spark.sql("select * from " + tablename).show()
```

id	first_name	last_name	email	gender	ip_address
1	Margaretta	Laughtisse	mlaughtisse0@medi...	Genderfluid	34.148.232.131
2	Vally	Garment	vgarment1@wisc.edu	Bigender	15.158.123.36
3	Tessa	Curee	tcuree2@php.net	Bigender	132.209.143.225
4	Arman	Heineking	aheineking3@tutto...	Male	157.110.61.233
5	Roselia	Trustie	rtrustie4@ft.com	Non-binary	49.55.218.81
6	Roxie	Springett	rspringett5@devia...	Male	51.206.104.138
7	Gabi	Kernell	gkernell6@hugedom...	Female	223.30.27.146
8	Dino	Kentwell	dkentwell7@com.com	Agender	107.244.52.181
9	Petronilla	Jandel	pjandel8@amazon.c...	Female	187.54.208.203
10	Courtney	Zecchinelli	czecchinelli9@cam...	Genderfluid	80.96.245.191
11	Sunny	Kennermann	skennermanna@quan...	Genderqueer	211.13.246.106
12	Dayle	McCrachen	dmccrachenb@booki...	Genderqueer	159.232.55.236
13	Cassie	Perschke	cperschkec@goo.gl	Genderqueer	193.62.46.4
14	Roshelle	Peskin	rpeskind@patch.com	Genderqueer	108.180.147.192
15	Carlie	Simonnot	csimonnote@survey...	Female	220.154.167.3
16	Gan	Siuda	gsiudaf@ucoz.ru	Agender	118.117.172.157
17	Klarika	Filimore	kfilimoreg@creati...	Agender	16.111.119.168
18	Lennie	Bilbrook	lbilbrookh@hatena...	Bigender	162.30.71.206
19	August	Cristoforo	acristoforo@goo...	Genderfluid	30.231.88.242
20	Avrit	Milburne	amilburnej@yahoo.com	Bigender	186.73.151.205

only showing top 20 rows

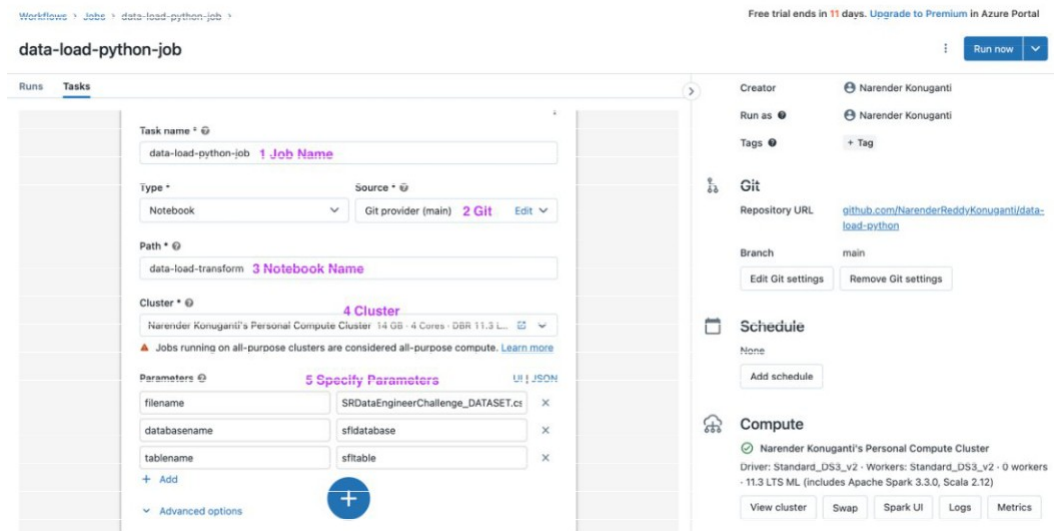
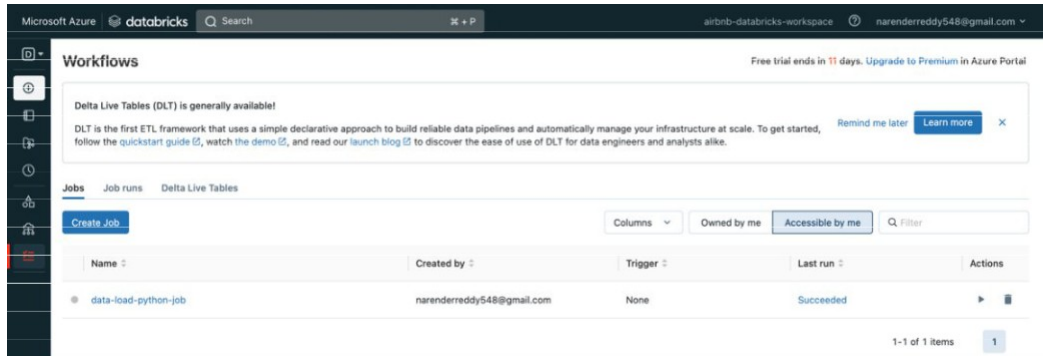
Conclusion

We have successfully extracted data from csv file, transformed and loaded it into the **database table** in databricks **data lakehouse**

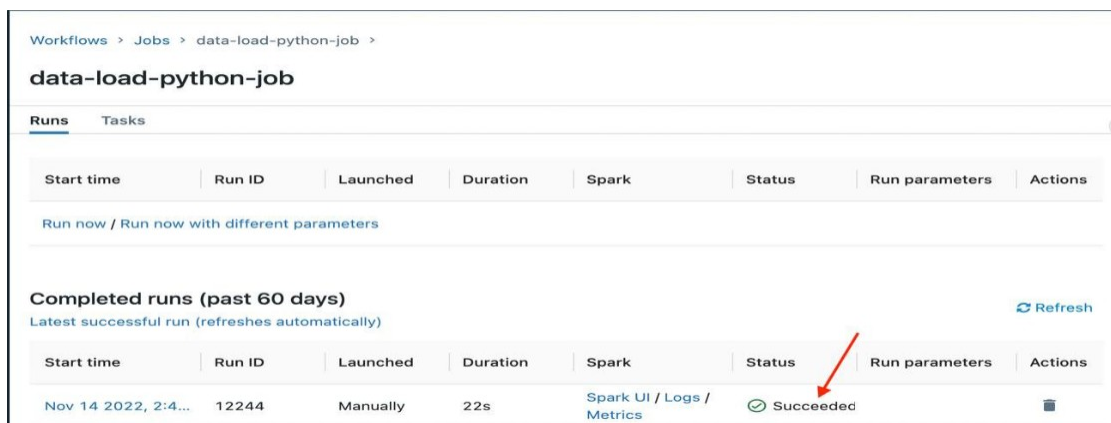
We can utilize the above data for running visualization and running ML models in Databricks.

Step 8: Create Job

- Σ Click on **Workflows** in the sidebar.
- Σ On the Workflows page, click Create Job. This opens the New Job page.
- Σ Specify a job name, git location, notebook name, Cluster name and parameters.



Job Run is successful



Step 9: Terraform configuration

Please follow the steps from [Configure Terraform Authentication](#) to *install requirements* and *Configure Terraform authentication* to authenticate the **Databricks Terraform provider** with your **Azure Databricks account** and your Azure Databricks workspace,

Follow [sample configuration](#) to provision a Databricks notebook, cluster, and a job to run the notebook on the cluster, in an existing Databricks workspace.

```
[(base) narenderreddykonuganti@Narenders-MacBook-Air terraform_cluster_notebook_job % pwd
/Users/narenderreddykonuganti/terraform_cluster_notebook_job
[(base) narenderreddykonuganti@Narenders-MacBook-Air terraform_cluster_notebook_job % ls -ltrh
total 88
-rw-r--r--  1 narenderreddykonuganti  staff   491B Nov 14 03:03 auth.tf
-rw-r--r--  1 narenderreddykonuganti  staff    42B Nov 14 03:03 auth.auto.tfvars
-rw-r--r--  1 narenderreddykonuganti  staff   1.1K Nov 14 03:05 cluster.tf
-rw-r--r--  1 narenderreddykonuganti  staff   656B Nov 14 03:05 notebook.tf
-rw-r--r--  1 narenderreddykonuganti  staff   137B Nov 14 03:07 notebook.auto.tfvars
-rw-r--r--  1 narenderreddykonuganti  staff   511B Nov 14 03:07 job.tf
-rw-r--r--  1 narenderreddykonuganti  staff    20B Nov 14 03:08 job.auto.tfvars
-rw-r--r--  1 narenderreddykonuganti  staff   3.5K Nov 14 03:26 data-load-transform.py
-rw-r--r--  1 narenderreddykonuganti  staff   3.9K Nov 14 03:27 terraform.tfstate.backup
-rw-r--r--  1 narenderreddykonuganti  staff   155B Nov 14 03:30 cluster.auto.tfvars
-rw-r--r--  1 narenderreddykonuganti  staff   3.9K Nov 14 03:34 terraform.tfstate
(base) narenderreddykonuganti@Narenders-MacBook-Air terraform_cluster_notebook_job %
```

me.tf --> This file gets information about the current databricks user

notebook.tf --> This file represents the notebook

notebook.auto.tfvars --> This file specifies the notebook's properties like notebook filename and language

data-load-transform.py --> This file represents the notebook's contents.

cluster.tf --> This file represents the cluster

cluster.auto.tfvars --> This file specifies the cluster's properties like cluster name and number of workers

job.tf --> This file represents the job that runs the notebook on the cluster.

job.auto.tfvars --> This file specifies the job's properties.

✓ Run terraform plan. If there are any errors, fix them, and then run the command again.

✓ **Run terraform apply**

When the above command is invoked, below are sequence of steps followed:

- Reads file **me.tf** and authenticates the user
- Reads **notebook.tf**, **notebook.auto.tfvars** and create notebook with content taken from **data-load-transform.py**
- Reads **cluster.tf**, **cluster.auto.tfvars** and creates cluster with specified name and no of workers
- Reads **job.tf**, **job.auto.tfvars** and creates job, assigns above created cluster to the job.

Verify that the notebook, cluster, and job were created: in the output of the terraform apply command, find the URLs for notebook_url, cluster_url, and job_url, and go to them.

Given the context of the use case, most important file is **data-load-transform.py** where the code with parameterization (*Filename, Database name, Table name*)

When you invoke the above created job with parameters as per environment, code will be executed.

Conclusion: Terraform code is reproducible to create **notebooks**, **clusters** and **jobs** which can be run with parameters based on environment

GitHub Repository Link:

<https://github.com/spappala1/ETLPipeline>